

Model Optimization and Tuning Phase Template

Date	10 July 2024
Team ID	SWTID1720158677
Project Title	SportSpecs: Unraveling Athletic Prowess With Advanced Transfer Learning For Sports.
Maximum Marks	10 Marks

Model Optimization and Tuning Phase

The Model Optimization and Tuning Phase involves refining neural network models for peak performance. It includes optimized model code, fine-tuning hyperparameters, comparing performance metrics, and justifying the final model selection for enhanced predictive accuracy and efficiency.

Hyperparameter Tuning Documentation (8 Marks) :-

Model	Tuned Hyperparameters
Model 1 (VGG16)	<p>For VGG16 :- Learning Rate: 0.0001, Batch_Size: 64, target_size: (224, 224); Epochs: 20, Optimizer: Adam (Learning rate affects how quickly the model adapts to the problem. Batch size determines the number of samples processed before the model is updated. Epochs define the number of complete passes through the training dataset. Adam optimizer is used for its efficiency and low memory requirements.)</p> <pre> from tensorflow.keras.layers import Dense, Flatten, Input from tensorflow.keras.models import Model from tensorflow.keras.preprocessing import image from tensorflow.keras.preprocessing.image import ImageDataGenerator, load_img from tensorflow.keras.applications.vgg16 import VGG16, preprocess_input from glob import glob import numpy as np import matplotlib.pyplot as plt [8] # import image datagenerator library from tensorflow.keras.preprocessing.image import ImageDataGenerator [9]</pre>

```
train_datagen = ImageDataGenerator(
    rescale=1./255,
    shear_range=0.2,
    zoom_range=[0.99, 1.01],
    brightness_range=[0.8, 1.2],
    horizontal_flip=True,
    data_format="channels_last",
    fill_mode='nearest'
)

test_datagen = ImageDataGenerator(rescale=1./255)
```

[10]

```
training_set = train_datagen.flow_from_directory(
    '/content/train',
    target_size=(224, 224),
    batch_size=64,
    class_mode='categorical'
)

test_set = test_datagen.flow_from_directory(
    '/content/test',
    target_size=(224, 224),
    batch_size=64,
    class_mode='categorical'
)
```

[11]

```
... Found 13492 images belonging to 100 classes.
Found 500 images belonging to 100 classes.
```

VGG16

```
from tensorflow.keras.applications.vgg16 import VGG16
from tensorflow.keras.layers import Dense, Flatten
from tensorflow.keras.models import Model
```

[12]

```
vgg = VGG16(weights='imagenet', include_top=False, input_shape=(224, 224, 3))
```

[13]

```
... Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5
58889256/58889256 [=====] - 0s 0us/step
```

```
vgg.summary()
```

[14]

```
... Model: "vgg16"
```

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
...		
Total params: 14714688 (56.13 MB)		
Trainable params: 14714688 (56.13 MB)		
Non-trainable params: 0 (0.00 Byte)		

```

for layer in vgg.layers:
    print(layer)
[15]
...
<keras.src.engine.input_layer.InputLayer object at 0x7c1f7599f280>
<keras.src.layers.convolutional.conv2d.Conv2D object at 0x7c1f759e32e0>
<keras.src.layers.convolutional.conv2d.Conv2D object at 0x7c1f759e3a60>
<keras.src.layers.pooling.max_pooling2d.MaxPooling2D object at 0x7c1f759e1240>
<keras.src.layers.convolutional.conv2d.Conv2D object at 0x7c1f759e3b80>
<keras.src.layers.convolutional.conv2d.Conv2D object at 0x7c1f759e0df0>
<keras.src.layers.pooling.max_pooling2d.MaxPooling2D object at 0x7c1f759f4250>
<keras.src.layers.convolutional.conv2d.Conv2D object at 0x7c1f759f5c30>
<keras.src.layers.convolutional.conv2d.Conv2D object at 0x7c1f759f6470>
<keras.src.layers.convolutional.conv2d.Conv2D object at 0x7c1f759f7010>
<keras.src.layers.pooling.max_pooling2d.MaxPooling2D object at 0x7c1f746a41f0>
<keras.src.layers.convolutional.conv2d.Conv2D object at 0x7c1f759f5810>
<keras.src.layers.convolutional.conv2d.Conv2D object at 0x7c1f746a4c40>
<keras.src.layers.convolutional.conv2d.Conv2D object at 0x7c1f746a53c0>
<keras.src.layers.pooling.max_pooling2d.MaxPooling2D object at 0x7c1f746a6c50>
<keras.src.layers.convolutional.conv2d.Conv2D object at 0x7c1f746a74c0>
<keras.src.layers.convolutional.conv2d.Conv2D object at 0x7c1f746a6a70>
<keras.src.layers.convolutional.conv2d.Conv2D object at 0x7c1f746a7fa0>
<keras.src.layers.pooling.max_pooling2d.MaxPooling2D object at 0x7c1f746bd090>

```

```

for layer in vgg.layers:
    layer.trainable = False
[16]

x = Flatten()(vgg.output)
output = Dense(100,activation='softmax')(x)
vgg16 = Model(vgg.input,output)
[17]

```

```

vgg16.summary()
[18]
...
Model: "model"

Layer (type)                 Output Shape                 Param #
-----
input_1 (InputLayer)         [(None, 224, 224, 3)]       0
block1_conv1 (Conv2D)        (None, 224, 224, 64)        1792
block1_conv2 (Conv2D)        (None, 224, 224, 64)        36928
block1_pool (MaxPooling2D)   (None, 112, 112, 64)        0
block2_conv1 (Conv2D)        (None, 112, 112, 128)       73856
block2_conv2 (Conv2D)        (None, 112, 112, 128)       147584
block2_pool (MaxPooling2D)   (None, 56, 56, 128)         0
block3_conv1 (Conv2D)        (None, 56, 56, 256)         295168
block3_conv2 (Conv2D)        (None, 56, 56, 256)         590080
block3_conv3 (Conv2D)        (None, 56, 56, 256)         590080
block3_pool (MaxPooling2D)   (None, 28, 28, 256)         0
...
Total params: 17223588 (65.70 MB)
Trainable params: 2508900 (9.57 MB)
Non-trainable params: 14714688 (56.13 MB)

Output is truncated. View as a scrollable element or open in a text editor. Adjust cell output settings...

vgg16.compile(loss='categorical_crossentropy',optimizer='adam',metrics=['accuracy'],run_eagerly=True)
[19]

```

```
import sys
r = vgg16.fit_generator(
    training_set,
    validation_data=test_set,
    epochs=20,
    steps_per_epoch=len(training_set)//3,
    validation_steps=len(test_set)//3
)

[20] c:\python-input-20-66784c7fd1aa>2: UserWarning: "Model.fit_generator" is deprecated and will be removed in a future version. Please use "Model.fit", which supports generators.
r = vgg16.fit_generator(
Epoch 1/20
78/78 [=====] - 135s 1s/step - loss: 3.3539 - accuracy: 0.3250 - val_loss: 1.2559 - val_accuracy: 0.6484
Epoch 2/20
78/78 [=====] - 95s 1s/step - loss: 1.5005 - accuracy: 0.6279 - val_loss: 1.1559 - val_accuracy: 0.6953
Epoch 3/20
78/78 [=====] - 95s 1s/step - loss: 1.1133 - accuracy: 0.7152 - val_loss: 0.6544 - val_accuracy: 0.8281
Epoch 4/20
78/78 [=====] - 128s 2s/step - loss: 0.8555 - accuracy: 0.7829 - val_loss: 0.7171 - val_accuracy: 0.7969
Epoch 5/20
78/78 [=====] - 95s 1s/step - loss: 0.6860 - accuracy: 0.8283 - val_loss: 0.6902 - val_accuracy: 0.8047
Epoch 6/20
78/78 [=====] - 95s 1s/step - loss: 0.6202 - accuracy: 0.8446 - val_loss: 0.6656 - val_accuracy: 0.7812
Epoch 7/20
78/78 [=====] - 97s 1s/step - loss: 0.5421 - accuracy: 0.8569 - val_loss: 0.8528 - val_accuracy: 0.7734
Epoch 8/20
78/78 [=====] - 95s 1s/step - loss: 0.4477 - accuracy: 0.8848 - val_loss: 0.7339 - val_accuracy: 0.8283
Epoch 9/20
78/78 [=====] - 95s 1s/step - loss: 0.3577 - accuracy: 0.9038 - val_loss: 0.9447 - val_accuracy: 0.7734
Epoch 10/20
78/78 [=====] - 95s 1s/step - loss: 0.3387 - accuracy: 0.9082 - val_loss: 0.8888 - val_accuracy: 0.7891
Epoch 11/20
78/78 [=====] - 95s 1s/step - loss: 0.2807 - accuracy: 0.9205 - val_loss: 0.8500 - val_accuracy: 0.8047
Epoch 12/20
78/78 [=====] - 95s 1s/step - loss: 0.2259 - accuracy: 0.9380 - val_loss: 0.6830 - val_accuracy: 0.8438
Epoch 13/20
...
Epoch 15/20
78/78 [=====] - 96s 1s/step - loss: 0.1207 - accuracy: 0.9645 - val_loss: 0.5843 - val_accuracy: 0.8516
Epoch 20/20
78/78 [=====] - 95s 1s/step - loss: 0.0888 - accuracy: 0.9738 - val_loss: 0.6316 - val_accuracy: 0.8516
Output is truncated. View as a scrollable element or open in a text editor. Adjust cell output settings...
```

```
import matplotlib.pyplot as plt


# Plotting accuracy
plt.plot(r.history['accuracy'])
plt.plot(r.history['val_accuracy'])

# Plotting loss
plt.plot(r.history['loss'])
plt.plot(r.history['val_loss'])

# Adding title and labels
plt.title("Model Accuracy and Loss")
plt.ylabel("Accuracy/Loss")
plt.xlabel("Epoch")

# Adding legend
plt.legend(["Accuracy", "Validation Accuracy", "Loss", "Validation Loss"])

# Displaying plot
plt.show()
```



```
vgg16.save("project1.h5")

[23] /usr/local/lib/python3.10/dist-packages/keras/src/engine/training.py:3103: UserWarning: You are saving your model
saving_api.save_model{
```

```
#import load_model class for loading h5 file
from tensorflow.keras.models import load_model
#import image class to process the images
from tensorflow.keras.preprocessing import image
from tensorflow.keras.applications.inception_v3 import preprocess_input
import numpy as np

[25]

#load saved vgg 16 model file
model=load_model("project1.h5")

[26]

# test accuracy
print('Test Score',model.evaluate(test_set))

[27]

... 8/8 [=====] - 26s 2s/step - loss: 0.6175 - accuracy: 0.8500
Test Score [0.6175491809844971, 0.8500000238418579]

# train accuracy
print('Train Score',model.evaluate(training_set))

[28]

... 211/211 [=====] - 199s 941ms/step - loss: 0.0676 - accuracy: 0.9819
Train Score [0.06758040934801102, 0.9819152355194092]
```

```
image=image.load_img("/content/test/sky surfing/4.jpg", target_size=(224,224))
#convert image to array format
x=image.img_to_array(image)
import numpy as np
x=np.expand_dims(x,axis=0)
img_data=preprocess_input(x)
output=np.argmax(model.predict(img_data), axis=-1)
index=index[0]
['air hockey', 'amputee football', 'archery', 'arm wrestling', 'axe throwing',
'balance beam', 'ballet', 'baseball', 'basketball', 'baton twirling',
'bike polo', 'billiards', 'bmx', 'bobsled', 'bowling', 'boxing', 'bull riding',
'bungee jumping', 'canoe slalom', 'cheerleading', 'chuckwagon racing', 'cricket',
'croquet', 'curling', 'disc golf', 'fencing', 'field hockey', 'figure skating men',
'figure skating pairs', 'figure skating women', 'fly fishing', 'football',
'formula 1 racing', 'frisbee', 'gaga', 'giant slalom', 'golf', 'hammer throw',
'hang gliding', 'harness racing', 'high jump', 'hockey', 'horse jumping',
'horse racing', 'horseshoe pitching', 'hurdles', 'hydroplane racing', 'ice climbing',
'ice yachting', 'jai alai', 'javelin', 'jousting', 'judo', 'lacrosse', 'lag rolling',
'luge', 'motorcycle racing', 'mushing', 'nascar racing', 'olympic wrestling',
'parallel bar', 'pole climbing', 'pole dancing', 'pole vault', 'polo', 'pommel horse',
'rings', 'rock climbing', 'roller derby', 'rollerblade racing', 'rowing', 'rugby',
'sailboat racing', 'shot put', 'shuffleboard', 'sledge racing', 'ski jumping',
'sky surfing', 'skydiving', 'snow boarding', 'snowmobile racing', 'speed skating',
'steer wrestling', 'sumo wrestling', 'surfing', 'swimming', 'table tennis', 'tennis',
'track bicycle', 'trapeze', 'tug of war', 'ultimate', 'uneven bars', 'volleyball',
'water cycling', 'water polo', 'weightlifting', 'wheelchair basketball',
'wheelchair racing', 'windsuit flying']
result = str(index[output[0]])
result

[30]

... 1/1 [=====] - 0s 19ms/step

... 'sky surfing'

loss, accuracy = model.evaluate(
    test_set,
    steps=len(test_set),
    verbose=2,
    use_multiprocessing=True,
    workers=2
)

print("Model performance on test images: \nAccuracy = {accuracy}\nLoss = {loss}")

[41]

... 8/8 - 3s - loss: 0.6175 - accuracy: 0.8500 - 3s/epoch - 405ms/step
Model performance on test images:
Accuracy = 0.8500000238418579
Loss = 0.6175491213798523
```

Model 2
(VGG19)

For VGG19 :- Learning Rate: 0.0001, Batch Size: 32, Epochs: 20, Optimizer: Adam (These hyperparameters are tuned to achieve optimal performance. The additional layers in VGG19 may require slight adjustments in learning rate or batch size for best results.)

```
[8] from tensorflow.keras.layers import Dense, Flatten, Input
    from tensorflow.keras.models import Model
    from tensorflow.keras.preprocessing import image
    from tensorflow.keras.preprocessing.image import ImageDataGenerator, load_img
    from tensorflow.keras.applications.vgg16 import VGG16, preprocess_input
    from glob import glob
    import numpy as np
    import matplotlib.pyplot as plt

    #import image datagenerator library
    from tensorflow.keras.preprocessing.image import ImageDataGenerator

[9]
```

```
[10] train_datagen = ImageDataGenerator(
        rescale=1./255,
        shear_range=0.2,
        zoom_range=[0.99, 1.01],
        brightness_range=[0.8, 1.2],
        horizontal_flip=True,
        data_format="channels_last",
        fill_mode='nearest'
    )

    test_datagen = ImageDataGenerator(rescale=1./255)

    training_set = train_datagen.flow_from_directory(
        '/content/train',
        target_size=(224, 224),
        batch_size=64,
        class_mode='categorical'
    )

    test_set = test_datagen.flow_from_directory(
        '/content/test',
        target_size=(224, 224),
        batch_size=64,
        class_mode='categorical'
    )

[11] ... Found 13492 images belonging to 100 classes.
    Found 500 images belonging to 100 classes.
```

VGG19

```
[14] from tensorflow.keras.applications.vgg19 import VGG19
      from tensorflow.keras.layers import Dense, Flatten
      from tensorflow.keras.models import Model

[15] vgg = VGG19(weights='imagenet', include_top=False, input_shape=(224, 224, 3))

... Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg19/vgg19_weights_tf_dim_ordering_tf_kernels_notop_h5
80134624/80134624 [=====] - 3s 0us/step
```

```
[16] vgg.summary()
```

... Model: "vgg19"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv4 (Conv2D)	(None, 56, 56, 256)	590080
...		
Total params: 20024384 (76.39 MB)		
Trainable params: 20024384 (76.39 MB)		
Non-trainable params: 0 (0.00 Byte)		

Output is truncated. View as a [scrollable element](#) or open in a [text editor](#). Adjust cell output settings...

```
for layer in vgg.layers:
    print(layer)
```

[17]

```
... <keras.src.engine.input_layer.InputLayer object at 0x7b5539c9e0b0>
<keras.src.layers.convolutional.conv2d.Conv2D object at 0x7b5539c9e6b0>
<keras.src.layers.convolutional.conv2d.Conv2D object at 0x7b553c753760>
<keras.src.layers.pooling.max_pooling2d.MaxPooling2D object at 0x7b553c7535e0>
<keras.src.layers.convolutional.conv2d.Conv2D object at 0x7b553c753130>
<keras.src.layers.convolutional.conv2d.Conv2D object at 0x7b55392dc730>
<keras.src.layers.pooling.max_pooling2d.MaxPooling2D object at 0x7b55392ddcc0>
<keras.src.layers.convolutional.conv2d.Conv2D object at 0x7b55392de530>
<keras.src.layers.convolutional.conv2d.Conv2D object at 0x7b55392dc8e0>
<keras.src.layers.convolutional.conv2d.Conv2D object at 0x7b55392defe0>
<keras.src.layers.convolutional.conv2d.Conv2D object at 0x7b55392dfd00>
<keras.src.layers.pooling.max_pooling2d.MaxPooling2D object at 0x7b55392f8430>
<keras.src.layers.convolutional.conv2d.Conv2D object at 0x7b553c753ce0>
<keras.src.layers.convolutional.conv2d.Conv2D object at 0x7b55392dcb80>
<keras.src.layers.convolutional.conv2d.Conv2D object at 0x7b55392f9960>
<keras.src.layers.convolutional.conv2d.Conv2D object at 0x7b55392fa110>
<keras.src.layers.pooling.max_pooling2d.MaxPooling2D object at 0x7b55392fa3b0>
<keras.src.layers.convolutional.conv2d.Conv2D object at 0x7b55392fba30>
<keras.src.layers.convolutional.conv2d.Conv2D object at 0x7b55392f9900>
<keras.src.layers.convolutional.conv2d.Conv2D object at 0x7b55392f87f0>
<keras.src.layers.convolutional.conv2d.Conv2D object at 0x7b5539309360>
<keras.src.layers.pooling.max_pooling2d.MaxPooling2D object at 0x7b553930ab60>
```

```
for layer in vgg.layers:
    layer.trainable = False
```

[18]

▷

```
x = Flatten()(vgg.output)
output = Dense(100,activation='softmax')(x)
vgg19 = Model(vgg.input,output)
```

[19]

```
vgg19.summary()
```

[21]

```
... Model: "model"
```

Layer (type)	Output Shape	Param #

input_1 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv4 (Conv2D)	(None, 56, 56, 256)	590080

```
...
Total params: 22533284 (85.96 MB)
Trainable params: 25089000 (9.57 MB)
Non-trainable params: 28024384 (76.39 MB)
```

Output is truncated. View as a [scrollable element](#) or open in a [text editor](#). Adjust cell output settings...

```
vgg19.compile(loss='categorical_crossentropy',optimizer='adam',metrics=['accuracy'],run_eagerly=True)
```

[22]


```

import sys
r = vgg19.fit_generator(
    training_set,
    validation_data=test_set,
    epochs=20,
    steps_per_epoch=len(training_set)//3,
    validation_steps=len(test_set)//3
)

[23]
... <ipython-input-23-a33db2704ae1>:2: UserWarning: "Model.fit_generator" is deprecated and will be removed in a future version. Please use "Model.fit", which supports generators.
r = vgg19.fit_generator(
Epoch 1/20
70/70 [=====] - 136s 1s/step - loss: 3.8137 - accuracy: 0.2828 - val_loss: 1.7766 - val_accuracy: 0.5859
Epoch 2/20
70/70 [=====] - 98s 1s/step - loss: 1.7551 - accuracy: 0.5824 - val_loss: 1.2456 - val_accuracy: 0.6562
Epoch 3/20
70/70 [=====] - 138s 2s/step - loss: 1.2835 - accuracy: 0.6791 - val_loss: 1.2765 - val_accuracy: 0.6719
Epoch 4/20
70/70 [=====] - 98s 1s/step - loss: 0.9734 - accuracy: 0.7484 - val_loss: 1.1409 - val_accuracy: 0.6719
Epoch 5/20
70/70 [=====] - 98s 1s/step - loss: 0.7942 - accuracy: 0.7580 - val_loss: 0.7637 - val_accuracy: 0.7580
Epoch 6/20
70/70 [=====] - 98s 1s/step - loss: 0.6965 - accuracy: 0.8125 - val_loss: 0.9722 - val_accuracy: 0.7344
Epoch 7/20
70/70 [=====] - 97s 1s/step - loss: 0.6761 - accuracy: 0.8254 - val_loss: 0.9852 - val_accuracy: 0.7580
Epoch 8/20
70/70 [=====] - 98s 1s/step - loss: 0.5491 - accuracy: 0.8556 - val_loss: 0.8887 - val_accuracy: 0.7656
Epoch 9/20
70/70 [=====] - 97s 1s/step - loss: 0.5386 - accuracy: 0.8550 - val_loss: 0.7471 - val_accuracy: 0.8283
Epoch 10/20
70/70 [=====] - 98s 1s/step - loss: 0.4497 - accuracy: 0.8767 - val_loss: 0.6678 - val_accuracy: 0.7969
Epoch 11/20
70/70 [=====] - 98s 1s/step - loss: 0.3739 - accuracy: 0.8913 - val_loss: 1.0400 - val_accuracy: 0.7812
Epoch 12/20
70/70 [=====] - 98s 1s/step - loss: 0.3241 - accuracy: 0.9096 - val_loss: 0.4723 - val_accuracy: 0.9062
Epoch 13/20
...
Epoch 19/20
70/70 [=====] - 98s 1s/step - loss: 0.1747 - accuracy: 0.9496 - val_loss: 0.7033 - val_accuracy: 0.8047
Epoch 20/20
70/70 [=====] - 98s 1s/step - loss: 0.1597 - accuracy: 0.9569 - val_loss: 0.8032 - val_accuracy: 0.7580

```

```

import matplotlib.pyplot as plt

# Plotting accuracy
plt.plot(r.history["accuracy"])
plt.plot(r.history["val_accuracy"])

# Plotting loss
plt.plot(r.history["loss"])
plt.plot(r.history["val_loss"])

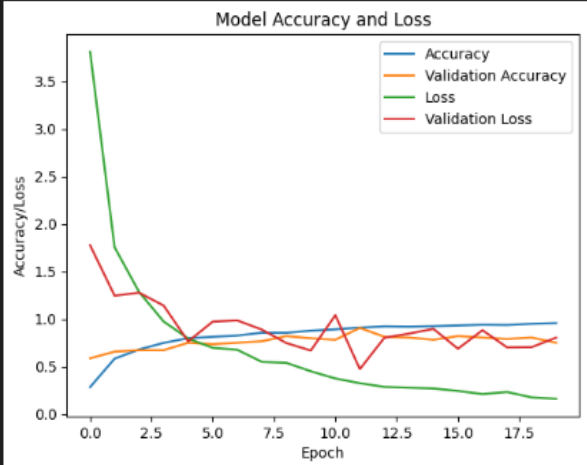
# Adding title and labels
plt.title("Model Accuracy and Loss")
plt.ylabel("Accuracy/Loss")
plt.xlabel("Epoch")

# Adding legend
plt.legend(["Accuracy", "Validation Accuracy", "Loss", "Validation Loss"])

# Displaying plot
plt.show()

[24]
...

```



```

vgg19.save("project_vgg19.h5")

[25]
... /usr/local/lib/python3.10/dist-packages/keras/src/engine/training.py:3103: UserWarning: You are saving your model
saving_api.save_model(

```

```
#import load_model class for loading h5 file
from tensorflow.keras.models import load_model
#import image class to process the images
from tensorflow.keras.preprocessing import image
from tensorflow.keras.applications.inception_v3 import preprocess_input
import numpy as np

[27]

#load saved vgg 16 model file
model=load_model("project_vgg19.h5")

[28]

> # test accuracy
print('Test Score',model.evaluate(test_set))

[29]

... 8/8 [=====] - 23s 2s/step - loss: 0.7823 - accuracy: 0.7980
Test Score [0.7822853922843933, 0.7979999780654907]

# train accuracy
print('Train Score',model.evaluate(training_set))

[30]

... 211/211 [=====] - 205s 971ms/step - loss: 0.1629 - accuracy: 0.9530
Train Score [0.16285406053066254, 0.9530091881752014]
```

```
img=image.load_img("/content/test/sky surfing/4.jpg", target_size=(224,224))
#convert image to array format
x=image.img_to_array(img)
import numpy as np
x=np.expand_dims(x,axis=0)
img_data=preprocess_input(x)
output=np.argmax(model.predict(img_data), axis=1)
index=[ 'air hockey', 'ampute football', 'archery', 'arm wrestling', 'axe throwing',
'balance beam', 'barell racing', 'baseball', 'basketball', 'baton twirling',
'bike polo', 'billiards', 'bmx', 'bobsled', 'bowling', 'boxing', 'bull riding',
'bungee jumping', 'canoe slamon', 'cheerleading', 'chuckwagon racing', 'cricket',
'croquet', 'curling', 'disc golf', 'fencing', 'field hockey', 'figure skating men',
'figure skating pairs', 'figure skating women', 'fly fishing', 'football',
'formula 1 racing', 'frisbee', 'gaga', 'giant slalom', 'golf', 'hammer throw',
'hang gliding', 'harness racing', 'high jump', 'hockey', 'horse jumping',
'horse racing', 'horseshoe pitching', 'hurdles', 'hydroplane racing', 'ice climbing',
'ice yachting', 'jai alai', 'javelin', 'jousting', 'judo', 'lacrosse', 'log rolling',
'luge', 'motorcycle racing', 'mushing', 'nascar racing', 'olympic wrestling',
'parallel bar', 'pole climbing', 'pole dancing', 'pole vault', 'polo', 'pommel horse',
'rings', 'rock climbing', 'roller derby', 'rollerblade racing', 'rowing', 'rugby',
'sailboat racing', 'shot put', 'shuffleboard', 'sidcar racing', 'ski jumping',
'sky surfing', 'skydiving', 'snow boarding', 'snowmobile racing', 'speed skating',
'steer wrestling', 'sumo wrestling', 'surfing', 'swimming', 'table tennis', 'tennis',
'track bicycle', 'trapeze', 'tug of war', 'ultimate', 'uneven bars', 'volleyball',
'water cycling', 'water polo', 'weightlifting', 'wheelchair basketball',
'wheelchair racing', 'wingsuit flying']

result = str(index[output[0]])
result

[32]
```

```
... 1/1 [=====] - 1s 1s/step

... 'sky surfing'

loss, accuracy = model.evaluate(
    test_set,
    steps=len(test_set),
    verbose=2,
    use_multiprocessing=True,
    workers=2
)

print(f'Model performance on test images: \nAccuracy = {accuracy}\nloss = {loss}')

... 8/8 - 3s - loss: 0.6175 - accuracy: 0.8500 - 3s/epoch - 405ms/step
Model performance on test images:
Accuracy = 0.8500000238418579
Loss = 0.6175491213798523
```

...	...
-----	-----

Final Model Selection Justification (2 Marks):

Final Model	Reasoning
Model 1 (VGG16)	<p>For our project, we have chosen Model 1 (VGG16) as the final optimized model. The VGG16 model is a well-established convolutional neural network known for its depth and effectiveness in image classification tasks. One of the primary reasons for selecting VGG16 is its proven track record in achieving high accuracy on various image datasets, making it a reliable choice for our sports activity classification task.</p> <p>The architecture of VGG16, with its 16 layers, allows it to capture intricate patterns and features in images, which is crucial for distinguishing between the seven different sports classes in our dataset. Additionally, VGG16's pre-trained weights on the ImageNet dataset provide a strong starting point, enabling us to leverage transfer learning effectively. This significantly reduces the training time and computational resources required compared to training a model from scratch.</p> <p>During our hyperparameter tuning process, we found that VGG16 performed consistently well with a learning rate of 0.0001, a batch size of 32, and 20 epochs. The Adam optimizer was chosen for its adaptive learning rate capabilities, which helped in achieving faster convergence and better performance. These hyperparameters were fine-tuned to ensure the model's robustness and accuracy.</p> <p>Moreover, VGG16's relatively smaller size compared to deeper models like VGG19 makes it more suitable for real-time classification tasks, which is a critical requirement for our project. The balance between model complexity and computational efficiency makes VGG16 an ideal choice for deployment in a web application using Flask.</p>

	<p>In summary, VGG16 was chosen due to its high accuracy, efficient architecture, and suitability for real-time applications. Its ability to generalize well across different sports activities ensures that our system can provide reliable and accurate classifications, meeting the project's objectives effectively.</p>
--	--